

TACO: CodeWeaver evidence profile

This profile audits every empirical surface identified for TACO and records the maximum scientifically defensible CodeWeaver comparison. Status: reference_only. 47.00% to 48.05% terminal-task accuracy is not project pass-all.

Evidence boundary

Paper values and CodeWeaver measurements remain separate unless subject identity, revision, denominator, and fixed oracle are verified. Missing values are not converted to zeros.

Existing result

results/crust-bench-codeweaver-comparison-2026-08-14

Complete empirical-surface audit

key	surface	denominator	metrics	status
taco	Table 1	6 terminal-agent benchmarks	baseline/TACO accuracy and token consumption	exact_reference_not_metric_aligned
taco	Table 2	TerminalBench 1.0/2.0 across about 9 models	terminal-agent performance	reference_only_not_metric_aligned
taco	Appendix ablations	terminal-agent benchmark runs	freeze rules, disable global evolution, best-of-K, compression overhead	reference_only_not_metric_aligned

Inclusion and execution decision

paper_id	source_url	empirical_scope	crust_role	codeweaver_status	reason	existing_result
arXiv:2604.19572	https://arxiv.org/abs/2604.19572	terminal-agent context compression over six benchmarks	CRUST-Bench terminal-task metric	reference_only	47.00% to 48.05% terminal-task accuracy is not project pass-all	results/crust-bench-codeweaver-comparison-2026-08-14

CRUST-Bench terminal-agent reference

backbone	baseline_accuracy_percent	taco_accuracy_percent	accuracy_delta_pp	baseline_tokens_thousands	taco_tokens_thousands	token_reduction_percent	denominator	metric
MiniMax-M2.5	47.0	48.05	1.05	163.53	134.97	17.5	not stated in Table 1	terminal-agent task accuracy