

CodeWeaver GPT-5.6 Sol comparison — interim snapshot

Captured: 2026-08-05T21:17:27.591104+00:00

Non-final: orchestration status only. Independent correctness evaluation and publication tables are pending the complete 354-cell matrix.

State	Cells
Completed	59
Genuine terminal failure	21
Running	53
Queued	221
Total	354

Protocol. 118 projects, three repetitions, GPT-5.6 Sol, maximum reasoning effort. Repetition 0 is preregistered for exact paper-table comparison; no best-of-three selection.

Baseline replay. Prior Oxidizer/AlphaTrans/SKEL artifacts are complete (18 measured). The 100 SWE-agent CRUST targets are explicitly unavailable and were not replaced with stubs. ReCodeAgent replay was still running when captured.

Integrity. Four telemetry-schema crashes were archived as infrastructure failures and excluded from measured outcomes. Every snapshot file is covered by metadata/checksums.sha256.

See README.md and metadata/interim_snapshot.json for the precise scope and review map.